

# Lecture 12: The IO Model & External Sorting

# Announcements

1. *Thank you for the great feedback (post coming soon)!*
2. Educational goals:
  1. ***Tech changes, principles change more slowly*** We teach principles and formal abstraction so you can adapt to a changing world and technology..
  2. **Ability to learn after you leave.** Why we give you new concepts in homeworks & projects. *We want you to be able to pick up those changing concepts. But we test you fairly.*
  3. **We select the essentials for you.** We've thought about the material quite a bit. Feedback helpful, but we'd hope to get the benefit of the doubt. 😊
3. Thank you for being awesome wrt the midterm.
  1. ... some of you started early... Not cool.
  2. SCPD people, a lot of you were *great*!

# Today's Lecture

1. *[From 9-2]*: Conflict Serializability & Deadlock
2. The Buffer
3. External Merge Sort

# 1. Conflict Serializability & Deadlock


*Recap from Lecture 9-2*

# What you will learn about in this section


1. RECAP: Concurrency
2. Conflict Serializability
3. DAGs & Topological Orderings
4. Strict 2PL
5. Deadlocks

# Recall: Concurrency as Interleaving TXNs

## Serial Schedule:



## Interleaved Schedule:



- For our purposes, having TXNs occur concurrently means **interleaving their component actions (R/W)**

We call the particular order of interleaving a schedule

# Recall: Why Interleave TXNs?

- Interleaving TXNs might lead to anomalous outcomes... why do it?
- Several important reasons:
  - Individual TXNs might be *slow*- don't want to block other users during!
  - Disk access may be *slow*- let some TXNs use CPUs while others accessing disk!

All concern large differences in *performance*

# Recall: Must Preserve Consistency & Isolation

- The DBMS has freedom to interleave TXNs
- However, it must pick an interleaving or **schedule** such that isolation and consistency are maintained
  - Must be *as if* the TXNs had executed serially!

“With great power comes great responsibility”


ACID

DBMS must pick a schedule which maintains isolation & consistency

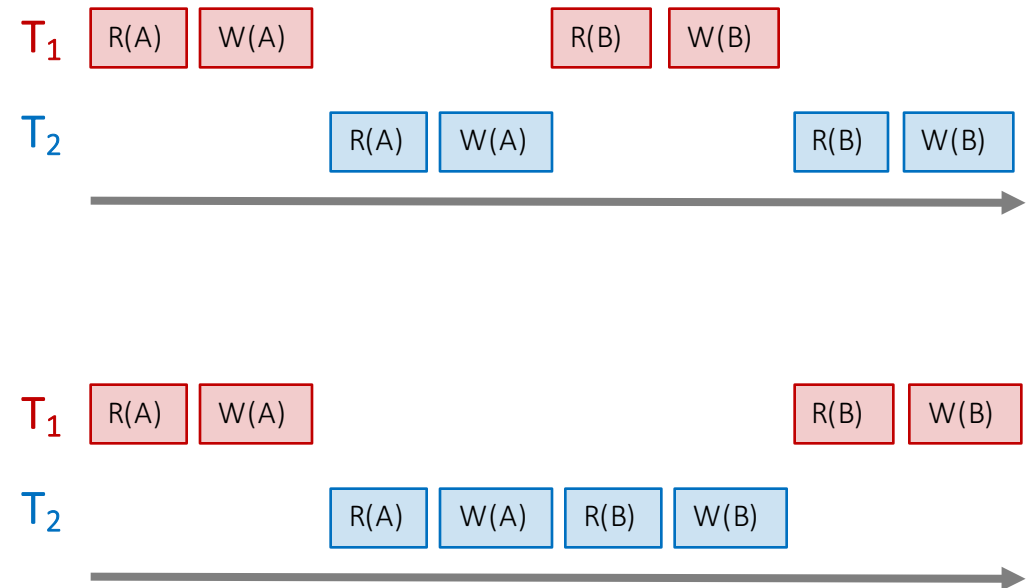


# Recall: “Good” vs. “bad” schedules

## Serial Schedule:



## Interleaved Schedules:



Why?

We want to develop ways of discerning “good” vs. “bad” schedules

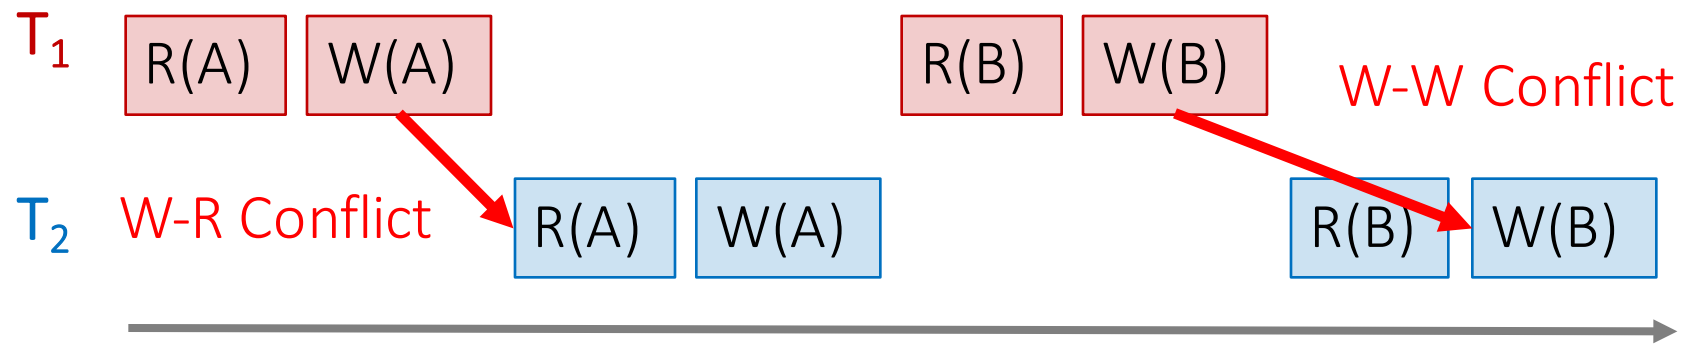
# Ways of Defining “Good” vs. “Bad” Schedules

- Recall from last time: we call a schedule ***serializable*** if it is equivalent to *some* serial schedule
  - We used this as a notion of a “good” interleaved schedule, since a **serializable schedule will maintain isolation & consistency**
- Now, we’ll define a stricter, but very useful variant:
  - **Conflict serializability**

We’ll need to define *conflicts* first..

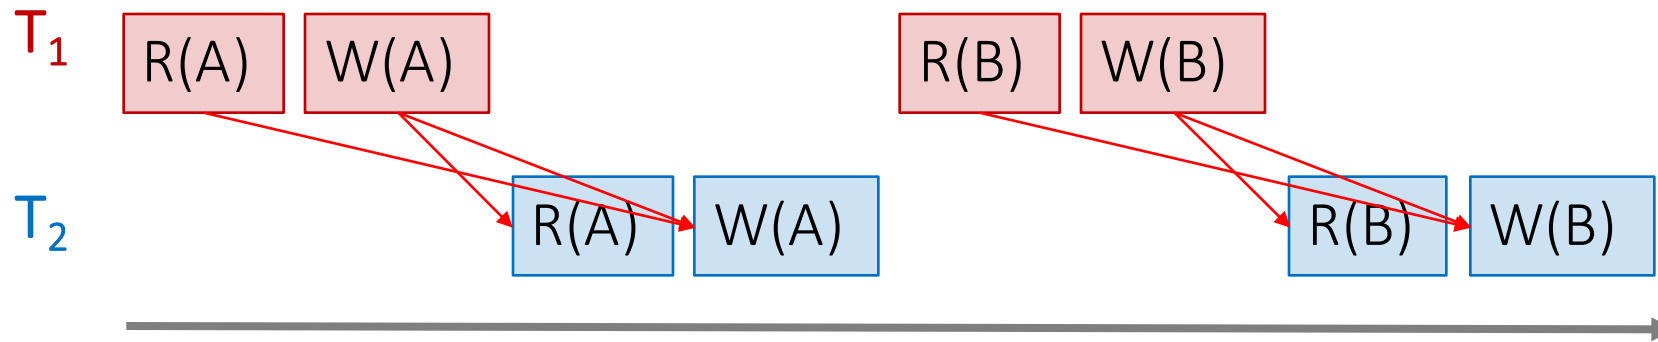
# Conflicts

Two actions conflict if they are part of different TXNs, involve the same variable, and at least one of them is a write



# Conflicts

Two actions conflict if they are part of different TXNs, involve the same variable, and at least one of them is a write



All “conflicts”!

# Conflict Serializability


- Two schedules are **conflict equivalent** if:
  - They involve *the same actions of the same TXNs*
  - Every *pair of conflicting actions* of two TXNs are *ordered in the same way*
- Schedule S is **conflict serializable** if S is *conflict equivalent* to some serial schedule

Conflict serializable  $\Rightarrow$  serializable

So if we have conflict serializable, we have consistency & isolation!


# Recall: “Good” vs. “bad” schedules

## Serial Schedule:



Note that in the “bad” schedule, the *order of conflicting actions is different than the above (or any) serial schedule!*

## Interleaved Schedules:




Conflict serializability also provides us with an operative notion of “good” vs. “bad” schedules!

# Note: Conflicts vs. Anomalies

- **Conflicts** are things we talk about to help us characterize different schedules
  - Present in both “good” and “bad” schedules
- **Anomalies** are instances where isolation and/or consistency is broken because of a “bad” schedule
  - We often characterize different anomaly types by what types of conflicts predicated them

# The Conflict Graph


- Let's now consider looking at conflicts **at the TXN level**
- Consider a graph where the **nodes are TXNs**, and there is an edge from  $T_i \rightarrow T_j$  if **any actions in  $T_i$  precede and conflict with any actions in  $T_j$**







# What can we say about “good” vs. “bad” conflict graphs?

## Serial Schedule:




A bit complicated...

## Interleaved Schedules:




# What can we say about “good” vs. “bad” conflict graphs?

Serial Schedule:



Simple!

Interleaved Schedules:



Theorem: Schedule is **conflict serializable** if and only if its conflict graph is acyclic


Let's unpack this notion of acyclic conflict graphs...

# DAGs & Topological Orderings

- A **topological ordering** of a directed graph is a linear ordering of its vertices that respects all the directed edges
- A directed **acyclic** graph (DAG) always has one or more **topological orderings**
  - (And there exists a topological ordering *if and only if* there are no directed cycles)

# DAGs & Topological Orderings


- Ex: What is one possible topological ordering here?



Ex: 0, 1, 2, 3 (or: 0, 1, 3, 2)

# DAGs & Topological Orderings

- Ex: What is one possible topological ordering here?



There is none!

# Connection to conflict serializability

- In the conflict graph, a topological ordering of nodes corresponds to a **serial ordering of TXNs**
- Thus an acyclic conflict graph  $\rightarrow$  conflict serializable!

Theorem: Schedule is **conflict serializable** if and only if its conflict graph is acyclic

# Strict Two-Phase Locking

- We consider **locking**- specifically, *strict two-phase locking*- as a way to deal with concurrency, because it **guarantees conflict serializability (if it completes- see upcoming...)**
- Also (*conceptually*) straightforward to implement, and transparent to the user!



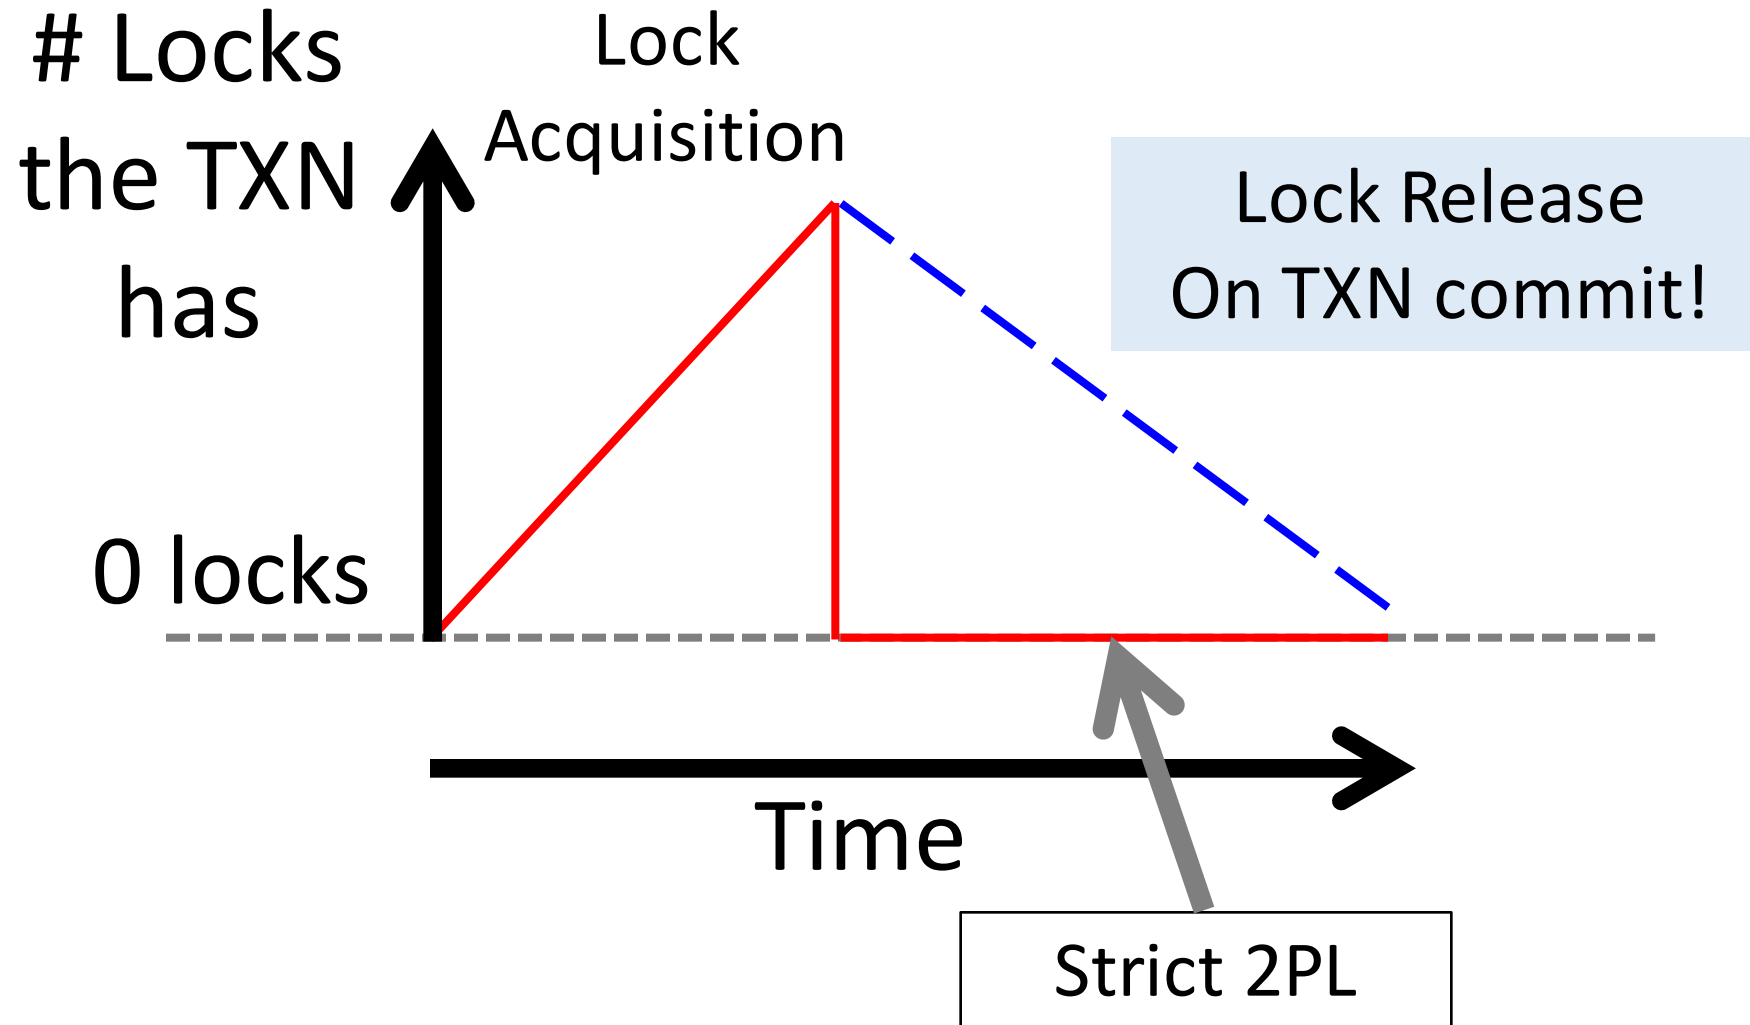
# Strict Two-phase Locking (Strict 2PL) Protocol:

## TXNs obtain:

- An **X (*exclusive*) lock** on object before **writing**.
  - If a TXN holds, no other TXN can get a lock (S or X) on that object.
- An **S (*shared*) lock** on object before **reading**
  - If a TXN holds, no other TXN can get an X lock on that object
- All locks held by a TXN are released when TXN completes.

Note: Terminology here- “exclusive”, “shared”- meant to be intuitive- no tricks!

# Picture of 2-Phase Locking (2PL)



# Strict 2PL

Theorem: Strict 2PL allows only schedules whose dependency graph is acyclic


*Proof Intuition*: In strict 2PL, if there is an edge  $T_i \rightarrow T_j$  (i.e.  $T_i$  and  $T_j$  conflict) then  $T_j$  needs to wait until  $T_i$  is finished – so *cannot* have an edge  $T_j \rightarrow T_i$

Therefore, Strict 2PL only allows conflict serializable  $\Rightarrow$  serializable schedules

# Strict 2PL

- If a schedule follows strict 2PL and locking, it is conflict serializable...
  - ...and thus serializable
  - ...and thus maintains isolation & consistency!
- Not all serializable schedules are allowed by strict 2PL.
- So let's use strict 2PL, what could go wrong?

# Deadlock Detection: Example



Waits-for graph:





First,  $T_1$  requests a shared lock on  $A$  to read from it

# Deadlock Detection: Example

$T_1$   $S(A)$   $R(A)$

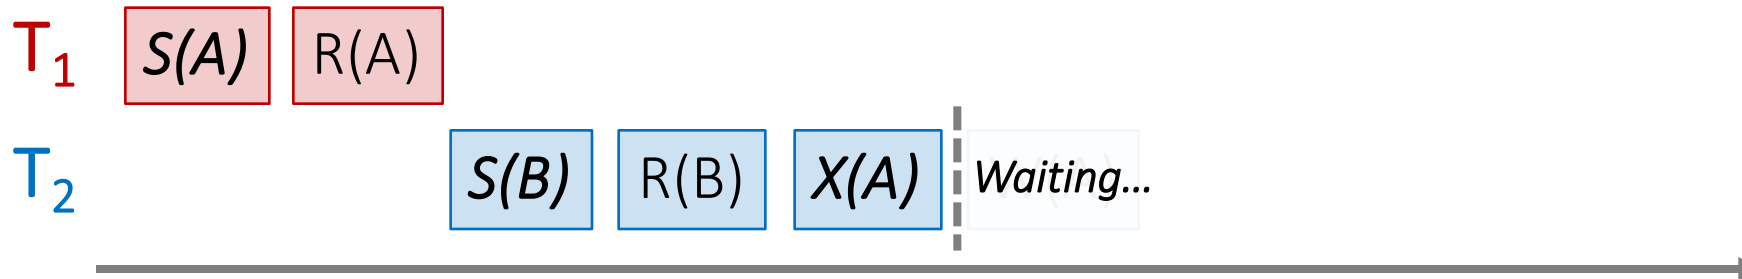
$T_2$   $S(B)$   $R(B)$

Waits-for graph:



Next,  $T_2$  requests a shared lock on B to read from it

# Deadlock Detection: Example

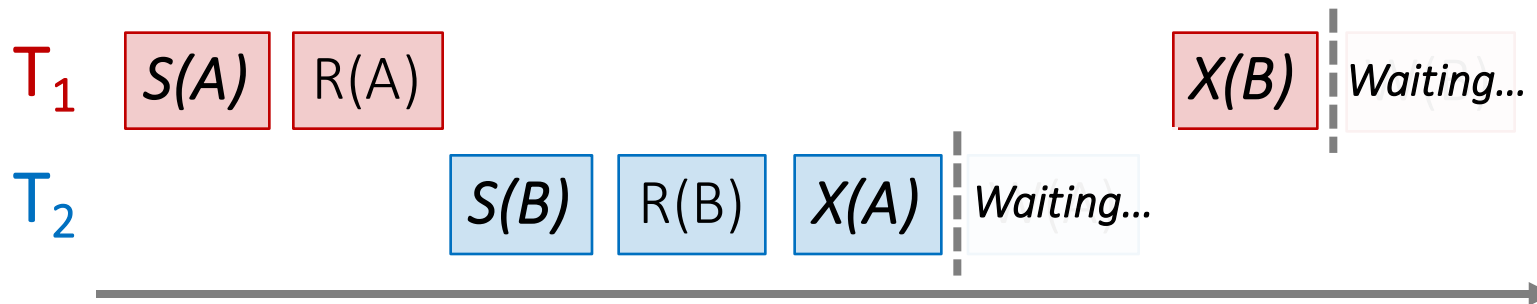


Waits-for graph:




$T_2$  then requests an exclusive lock on  $A$  to write to it- **now**  $T_2$  is waiting on  $T_1$ ...

# Deadlock Detection: Example



Finally,  $T_1$  requests an exclusive lock on B to write to it- **now  $T_1$  is waiting on  $T_2$ ... DEADLOCK!**

Waits-for graph:




Cycle =  
DEADLOCK




```
sqlite3.OperationalError: database is locked
```

```
ERROR: deadlock detected
DETAIL: Process 321 waits for ExclusiveLock on tuple of
relation 20 of database 12002; blocked by process 4924.
Process 404 waits for ShareLock on transaction 689; blocked
by process 552.
HINT: See server log for query details.
```



The problem?  
Deadlock!??!



NB: Also movie called wedlock  
(deadlock) set in a futuristic prison...  
I haven't seen either of them...

# Deadlocks

- **Deadlock:** Cycle of transactions waiting for locks to be released by each other.
- Two ways of dealing with deadlocks:
  1. Deadlock prevention
  2. Deadlock detection

# Deadlock Detection

- Create the **waits-for graph**:
  - Nodes are transactions
  - There is an edge from  $T_i \rightarrow T_j$  if  $T_i$  is *waiting for  $T_j$  to release a lock*
- Periodically check for (***and break***) cycles in the waits-for graph

# Summary

- *Last lecture:* Concurrency achieved by **interleaving TXNs** such that **isolation & consistency** are maintained
  - We formalized a notion of **serializability** that captured such a “good” interleavingschedule
- We defined **conflict serializability**, which implies serializability
  - There are other, more general issues!
- **Locking** allows only conflict serializable schedules
  - If the schedule completes- it may deadlock!



Candy Break

## 2. The Buffer

# Transition to Mechanisms


1. So you can **understand** what the database is doing!
  1. Understand the CS challenges of a database and how to use it.
  2. Understand how to optimize a query
  
2. Many **mechanisms** have become **stand-alone systems**
  - **Indexing** to Key-value stores
  - Embedded join processing
  - SQL-like languages take some aspect of what we discuss (PIG, Hive)

# What you will learn about in this section

1. RECAP: Storage and memory model
2. Buffer primer



# High-level: Disk vs. Main Memory



## Disk:


- **Slow:** Sequential *block* access
  - Read a blocks (not byte) at a time, so sequential access is cheaper than random
  - **Disk read / writes are expensive!**
- **Durable:** We will assume that once on disk, data is safe!
- **Cheap**

## Random Access Memory (RAM) or Main Memory:

- **Fast:** Random access, byte addressable
  - ~10x faster for sequential access
  - ~100,000x faster for random access!
- **Volatile:** Data can be lost if e.g. crash occurs, power goes out, etc!
- **Expensive:** For \$100, get 16GB of RAM vs. 2TB of disk!


# The Buffer

- A **buffer** is a region of physical memory used to store *temporary data*
  - *In this lecture:* a region in main memory used to store **intermediate data between disk and processes**
- *Key idea:* Reading / writing to disk is slow-need to cache data!



# The (Simplified) Buffer


- In this class: We'll consider a buffer located in **main memory** that operates over **pages** and **files**:
  - **Read(page)**: Read page from disk -> buffer *if not already in buffer*



# The (Simplified) Buffer


- In this class: We'll consider a buffer located in **main memory** that operates over **pages** and **files**:
- **Read(page)**: Read page from disk -> buffer *if not already in buffer*

Processes can then read from / write to the page in the buffer




# The (Simplified) Buffer

- In this class: We'll consider a buffer located in **main memory** that operates over **pages** and **files**:
  - **Read(page)**: Read page from disk -> buffer *if not already in buffer*
  - **Flush(page)**: Evict page from buffer & write to disk




# The (Simplified) Buffer

- In this class: We'll consider a buffer located in **main memory** that operates over **pages** and **files**:
  - **Read(page)**: Read page from disk -> buffer *if not already in buffer*
  - **Flush(page)**: Evict page from buffer & write to disk
  - **Release(page)**: Evict page from buffer *without* writing to disk



# Managing Disk: The DBMS Buffer

- Database maintains its own buffer
  - Why? The OS already does this...
  - DB knows more about access patterns.
    - Watch for how this shows up! (cf. *Sequential Flooding*)
  - Recovery and logging require ability to **flush** to disk.




# The Buffer Manager

- A **buffer manager** handles supporting operations for the buffer:
  - Primarily, handles & executes the “replacement policy”
    - i.e. finds a page in buffer to flush/release if buffer is full and a new page needs to be read in
- DBMSs typically implement their own buffer management routines



# A Simplified Filesystem Model

- For us, a **page** is a ***fixed-sized array*** of memory
  - Think: One or more disk blocks
  - Interface:
    - write to an entry (called a **slot**) or set to “None”
  - DBMS also needs to handle variable length fields
    - Page layout is important for good hardware utilization as well (see 346)
- And a **file** is a *variable-length list* of pages
  - Interface: create / open / close; next\_page(); etc.



## 2. External Merge & Sort

# What you will learn about in this section

1. External Merge- Basics
2. External Merge- Extensions
3. External Sort

# External Merge

# Challenge: Merging Big Files with Small Memory

How do we *efficiently* merge two sorted files when both are much larger than our main memory buffer?

# External Merge Algorithm

- **Input:** 2 sorted lists of length  $M$  and  $N$
- **Output:** 1 sorted list of length  $M + N$
- **Required:** At least 3 Buffer Pages
- **IOs:**  $2(M+N)$

# Key (Simple) Idea

To find an element that is no larger than all elements in two lists, one only needs to compare minimum elements from each list.

If:

$$A_1 \leq A_2 \leq \dots \leq A_N$$

$$B_1 \leq B_2 \leq \dots \leq B_M$$

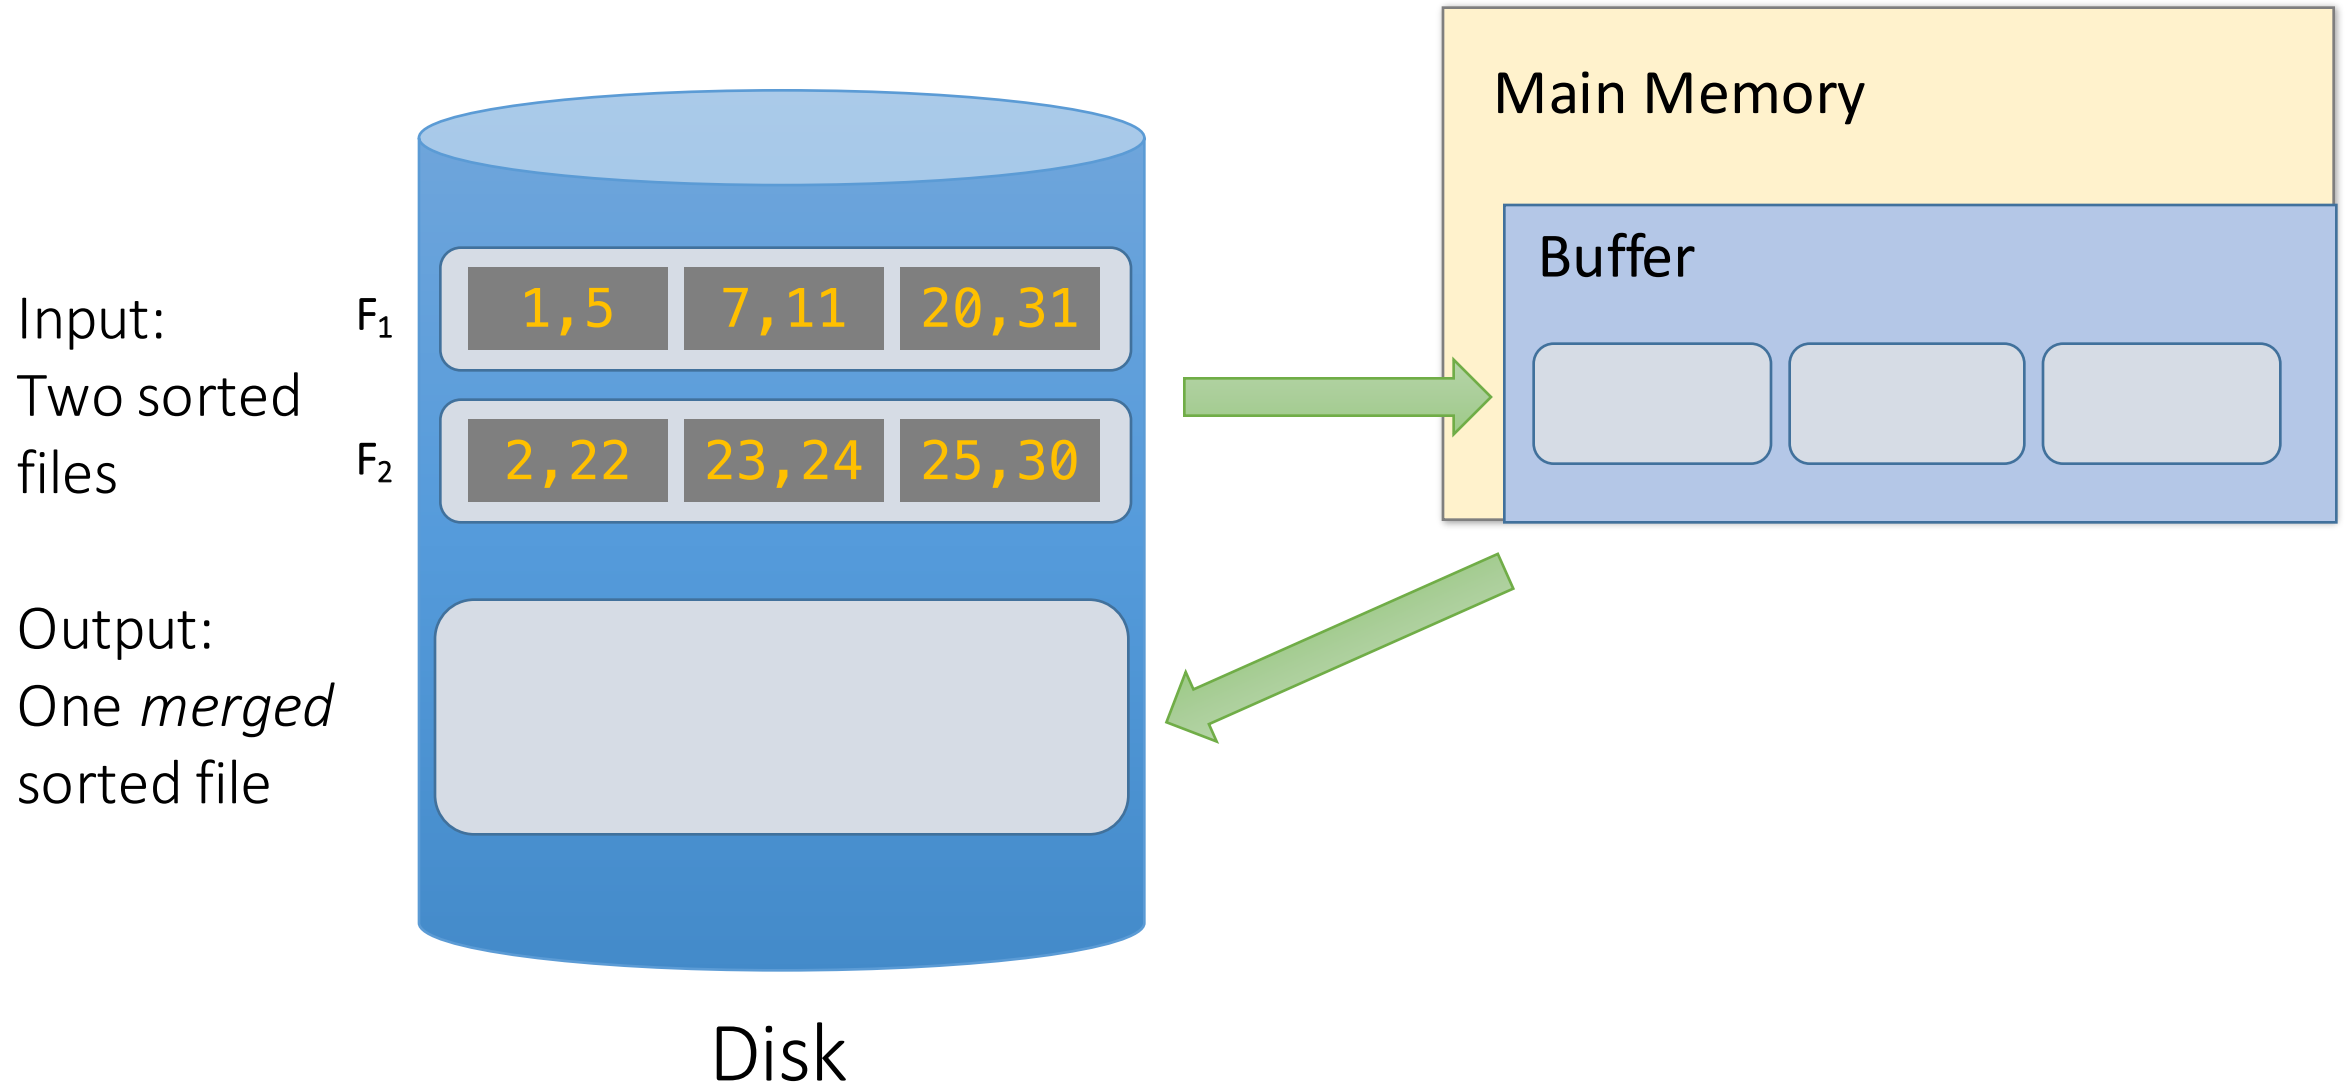
Then:

$$\text{Min}(A_1, B_1) \leq A_i$$

$$\text{Min}(A_1, B_1) \leq B_j$$

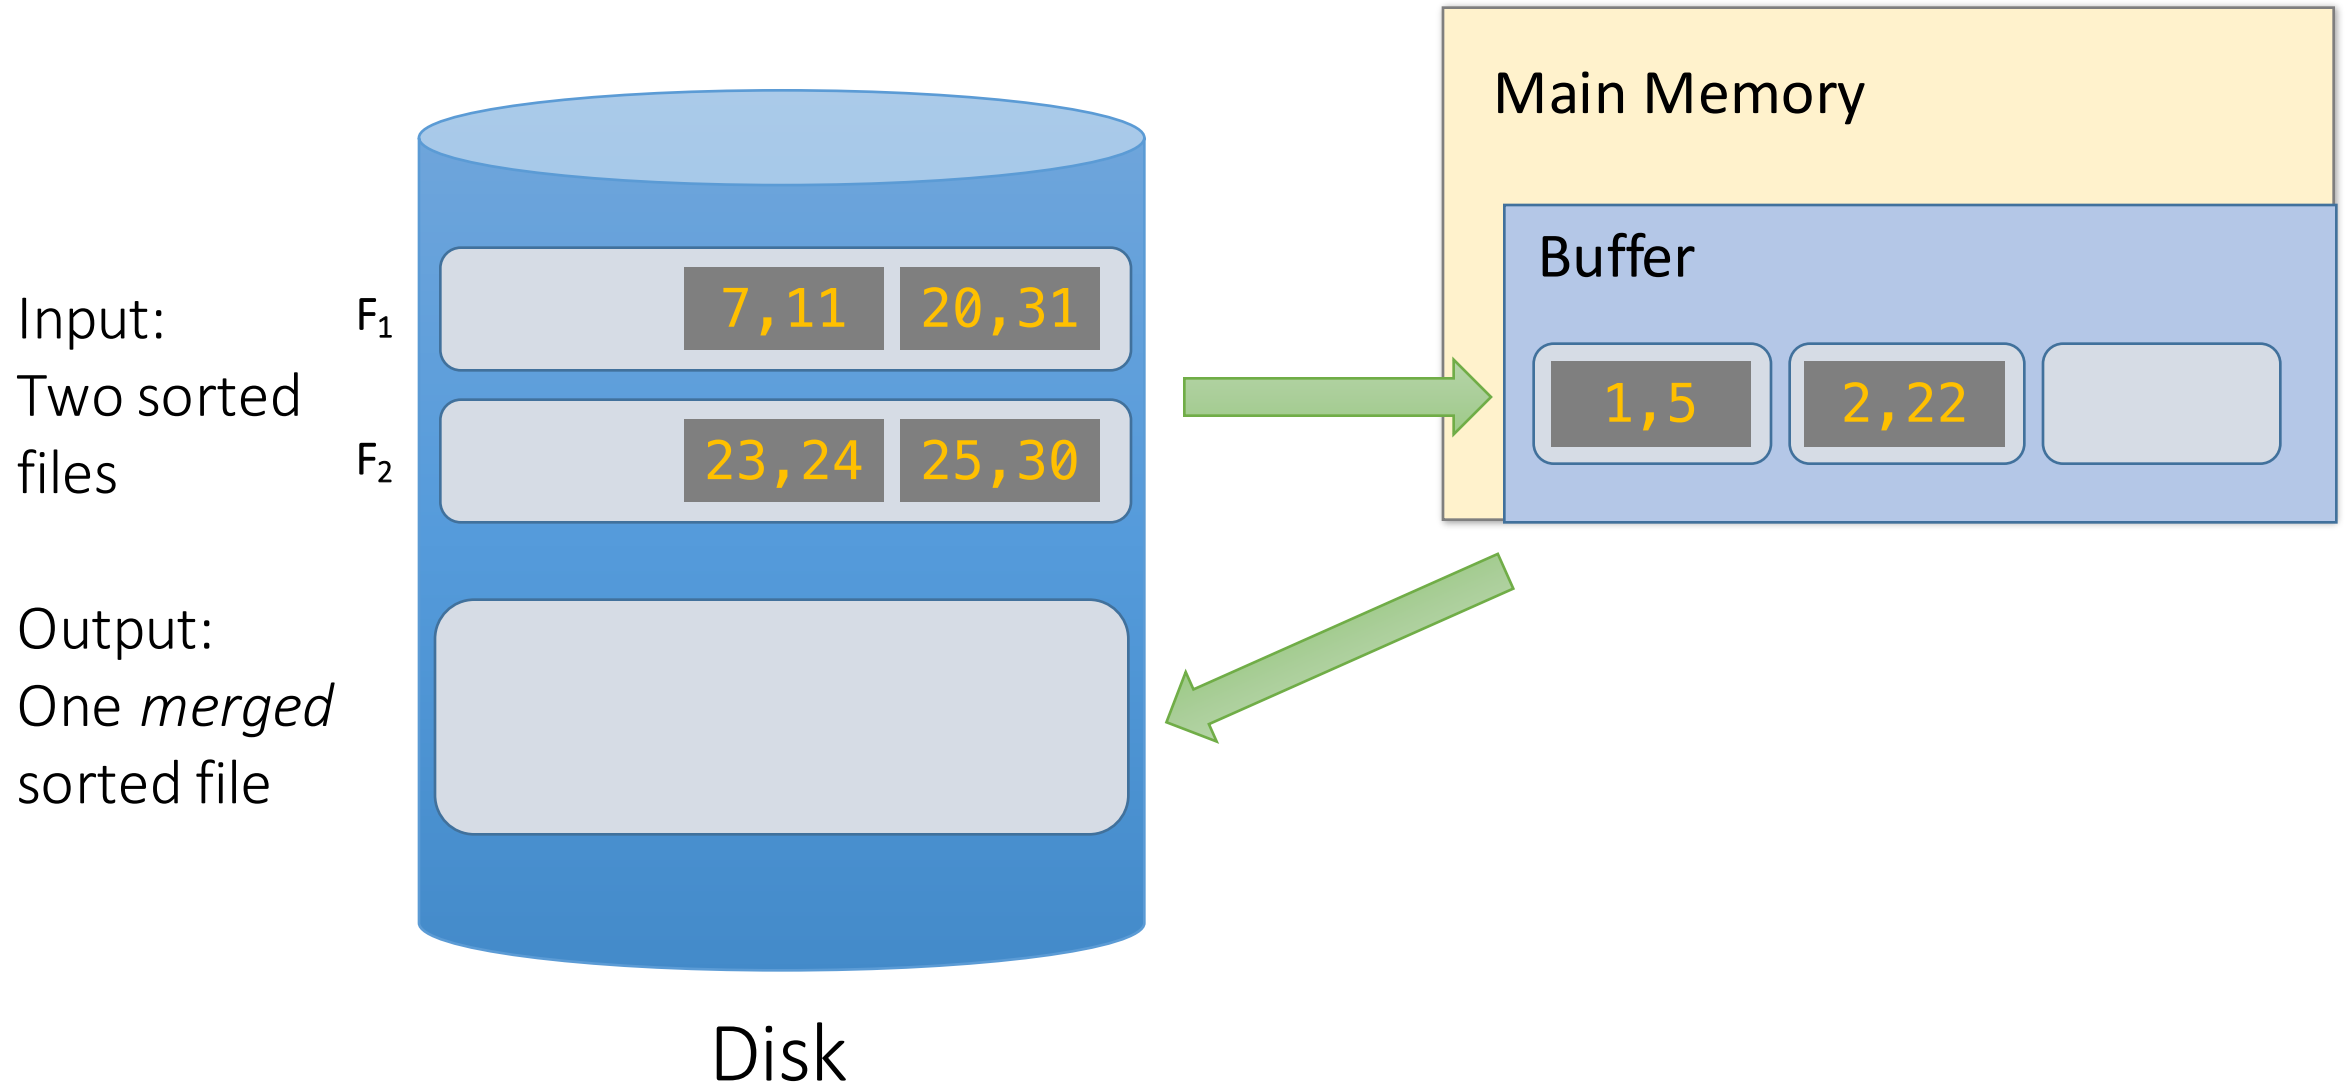
for  $i=1\dots N$  and  $j=1\dots M$

# External Merge Algorithm

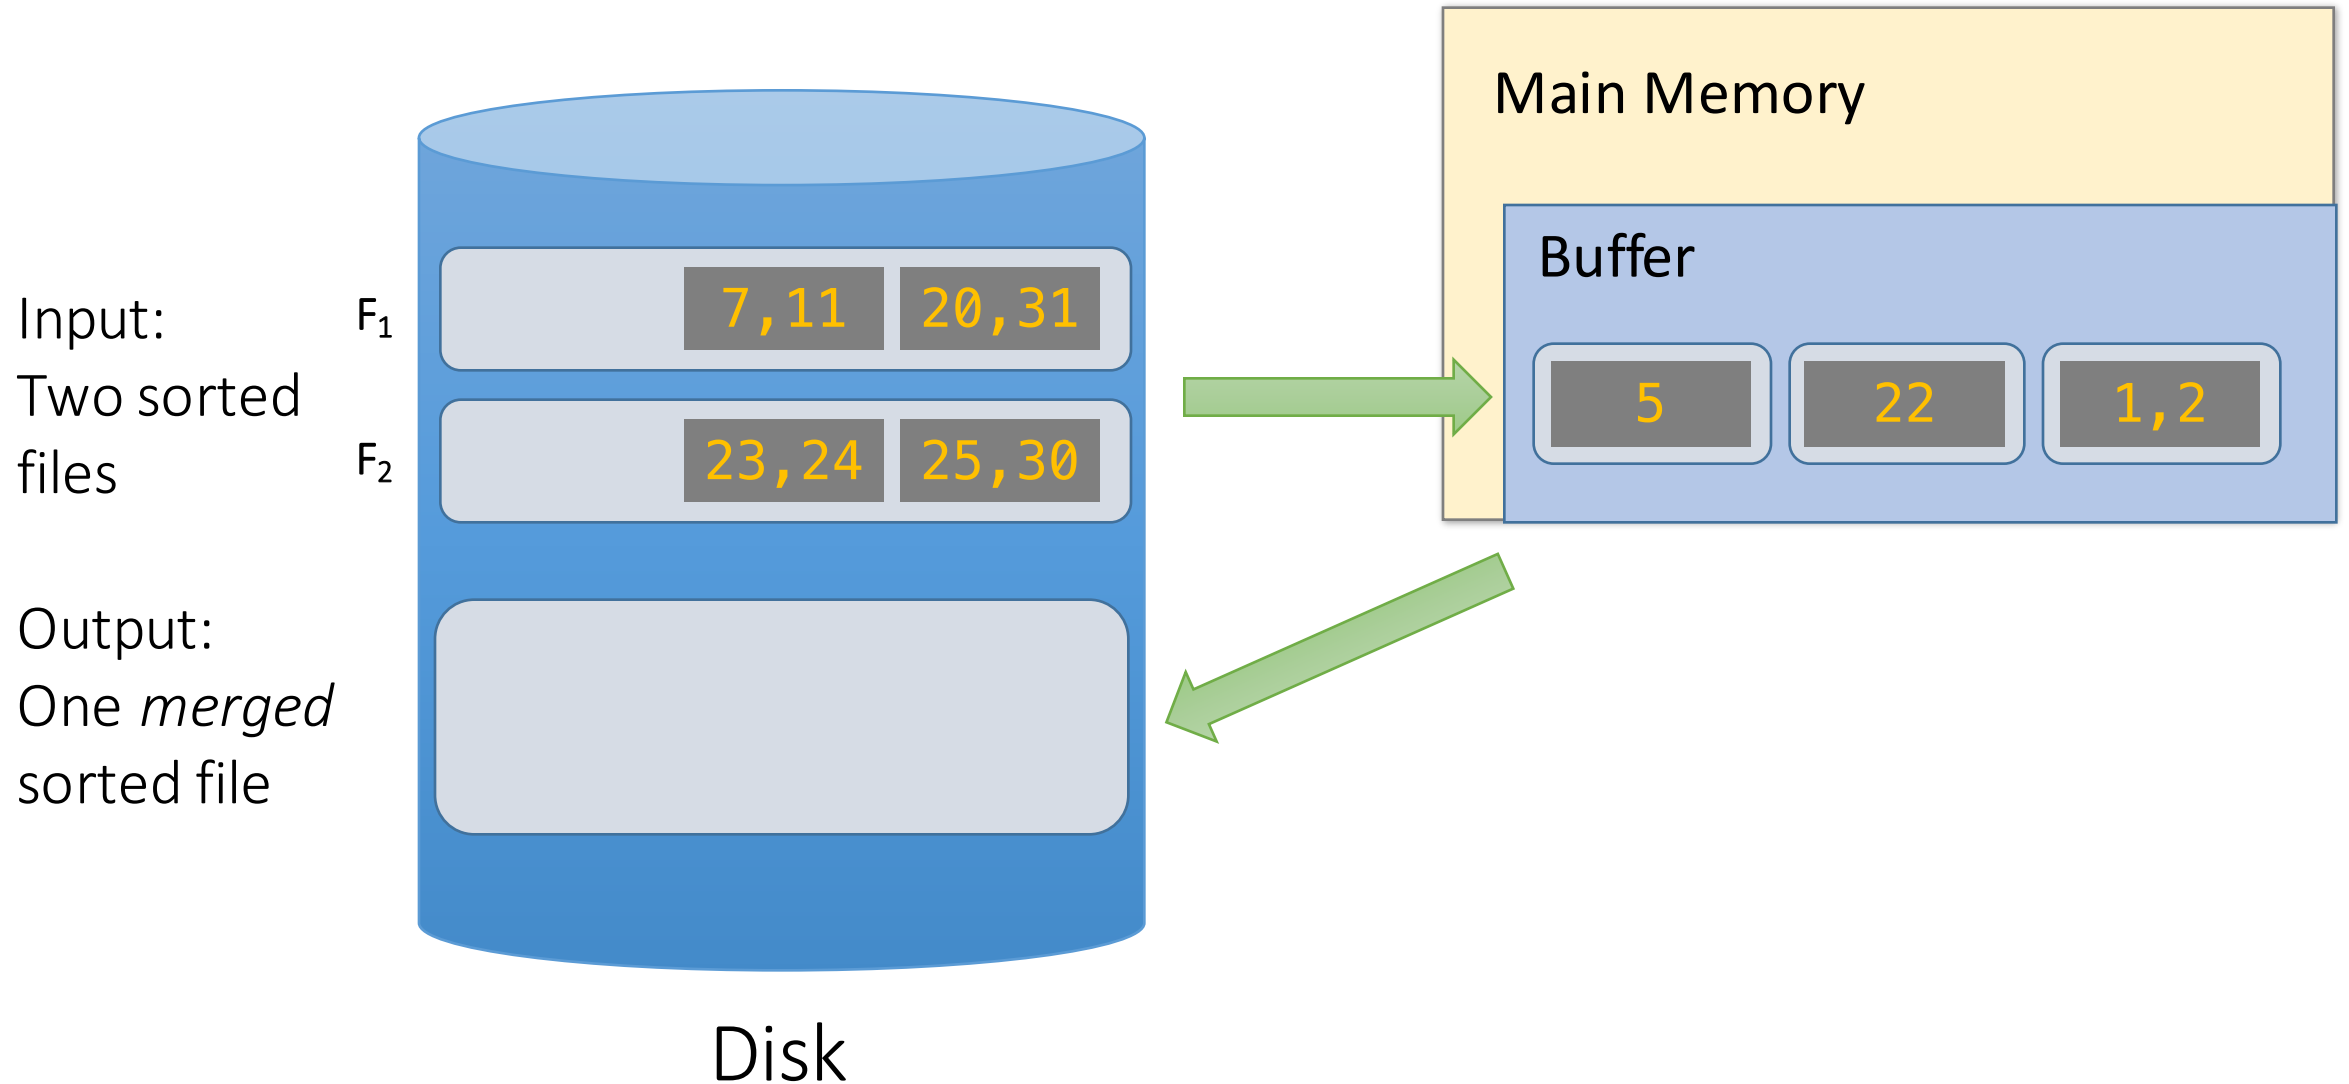




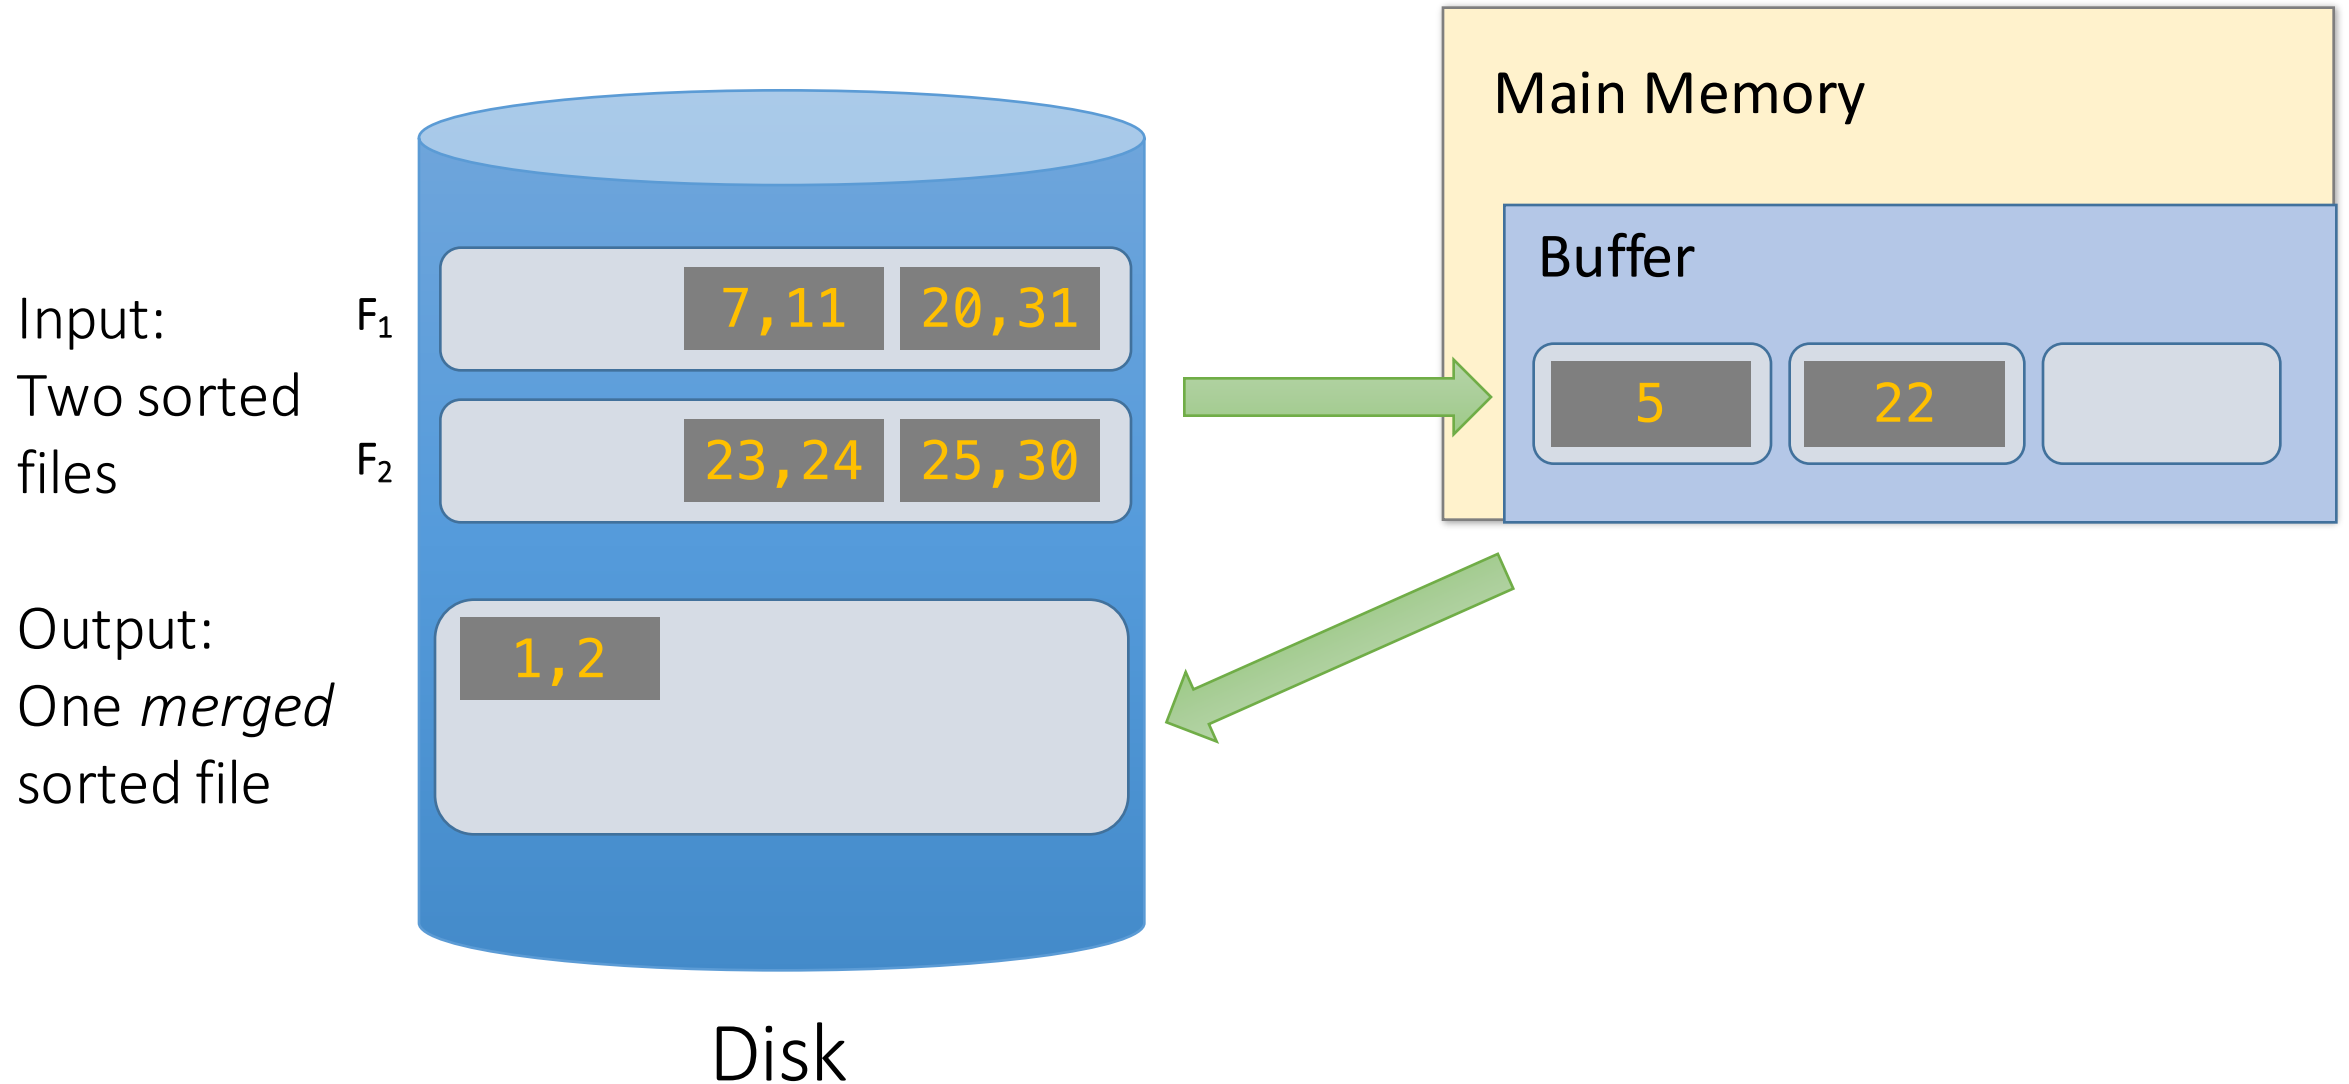
# External Merge Algorithm



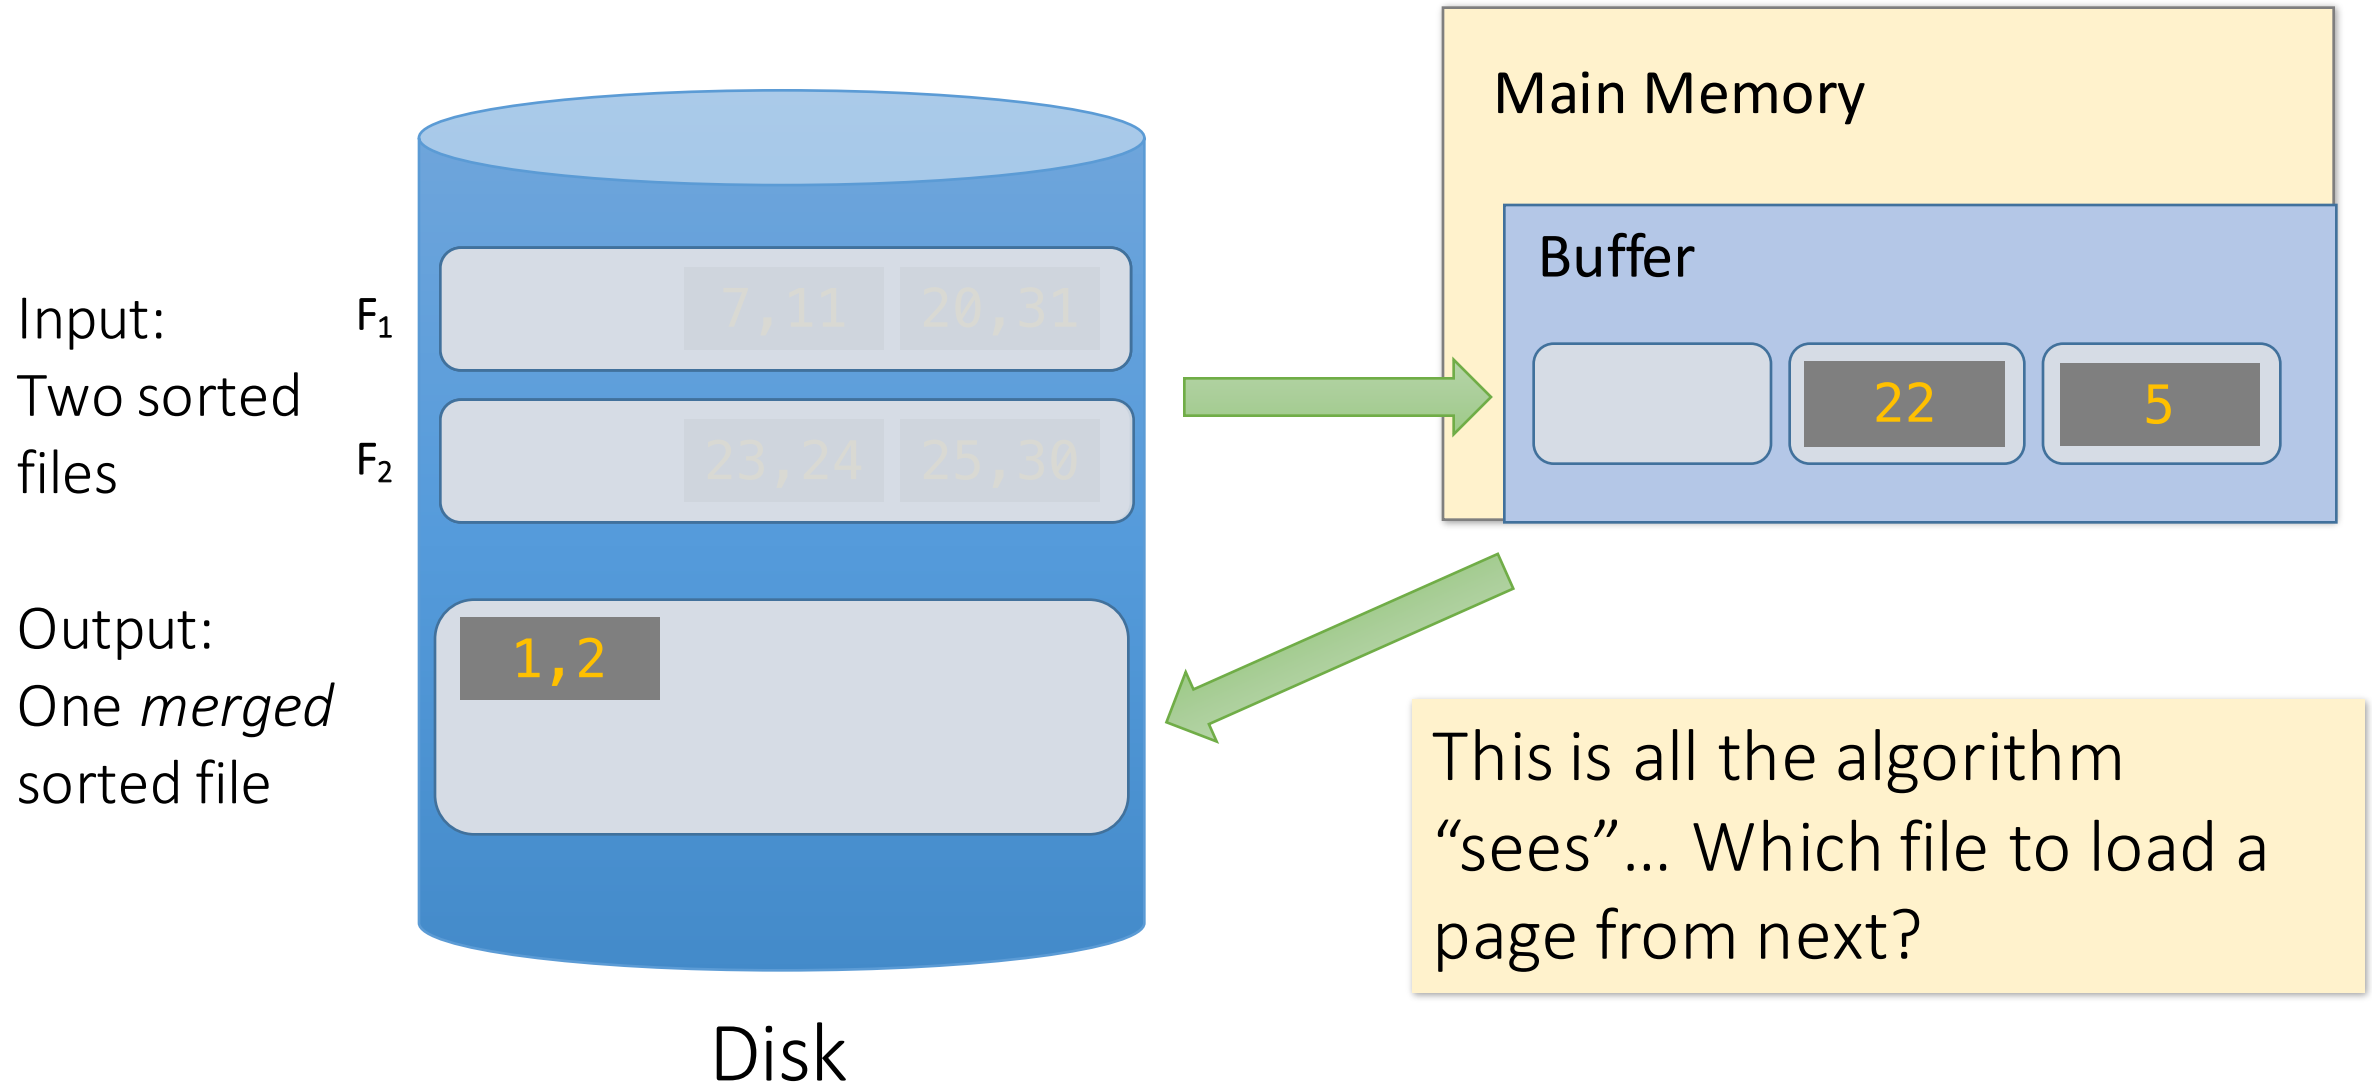
# External Merge Algorithm



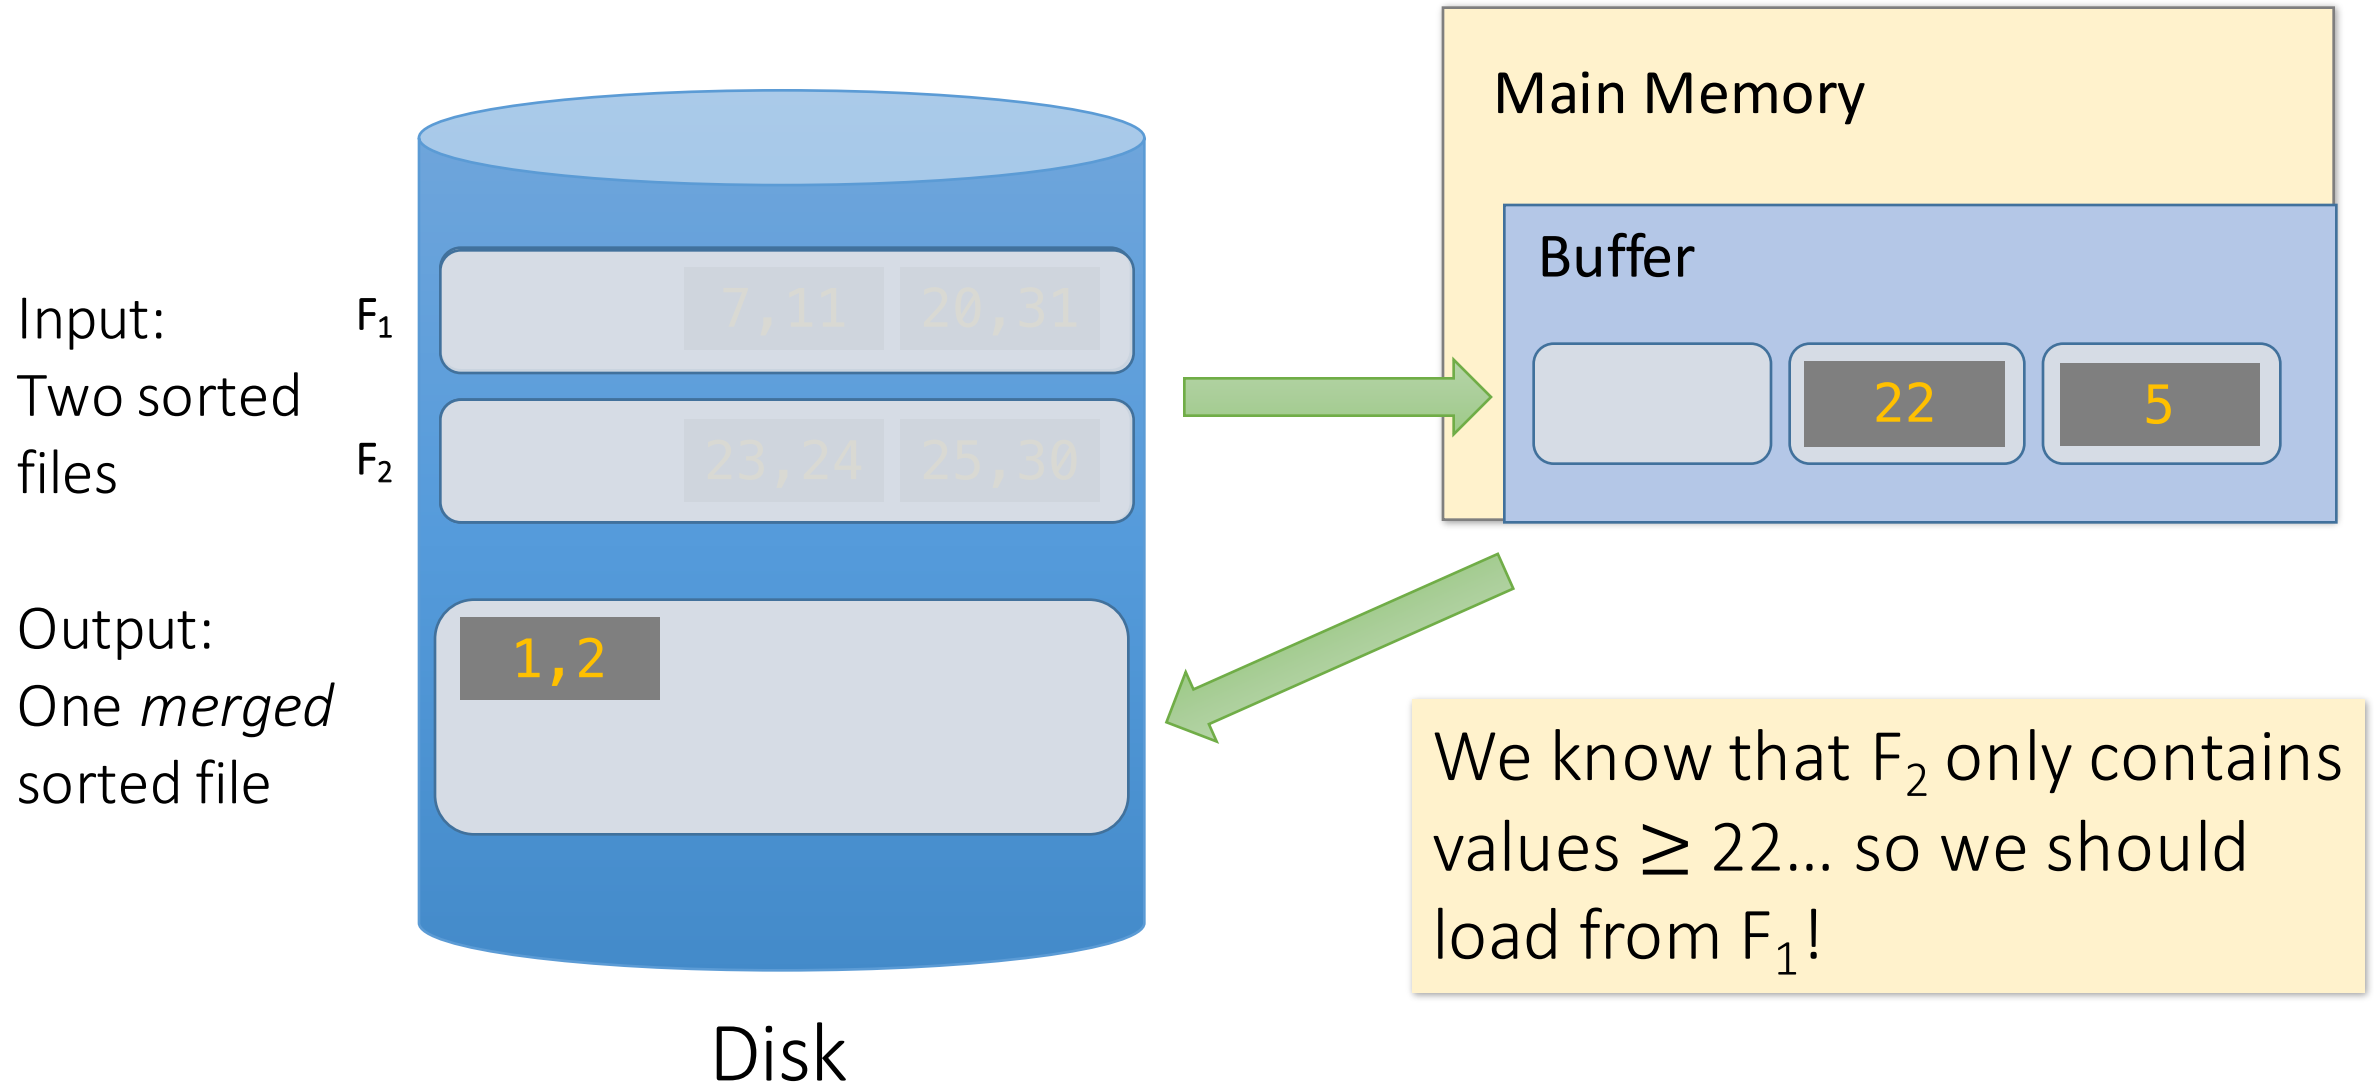
# External Merge Algorithm



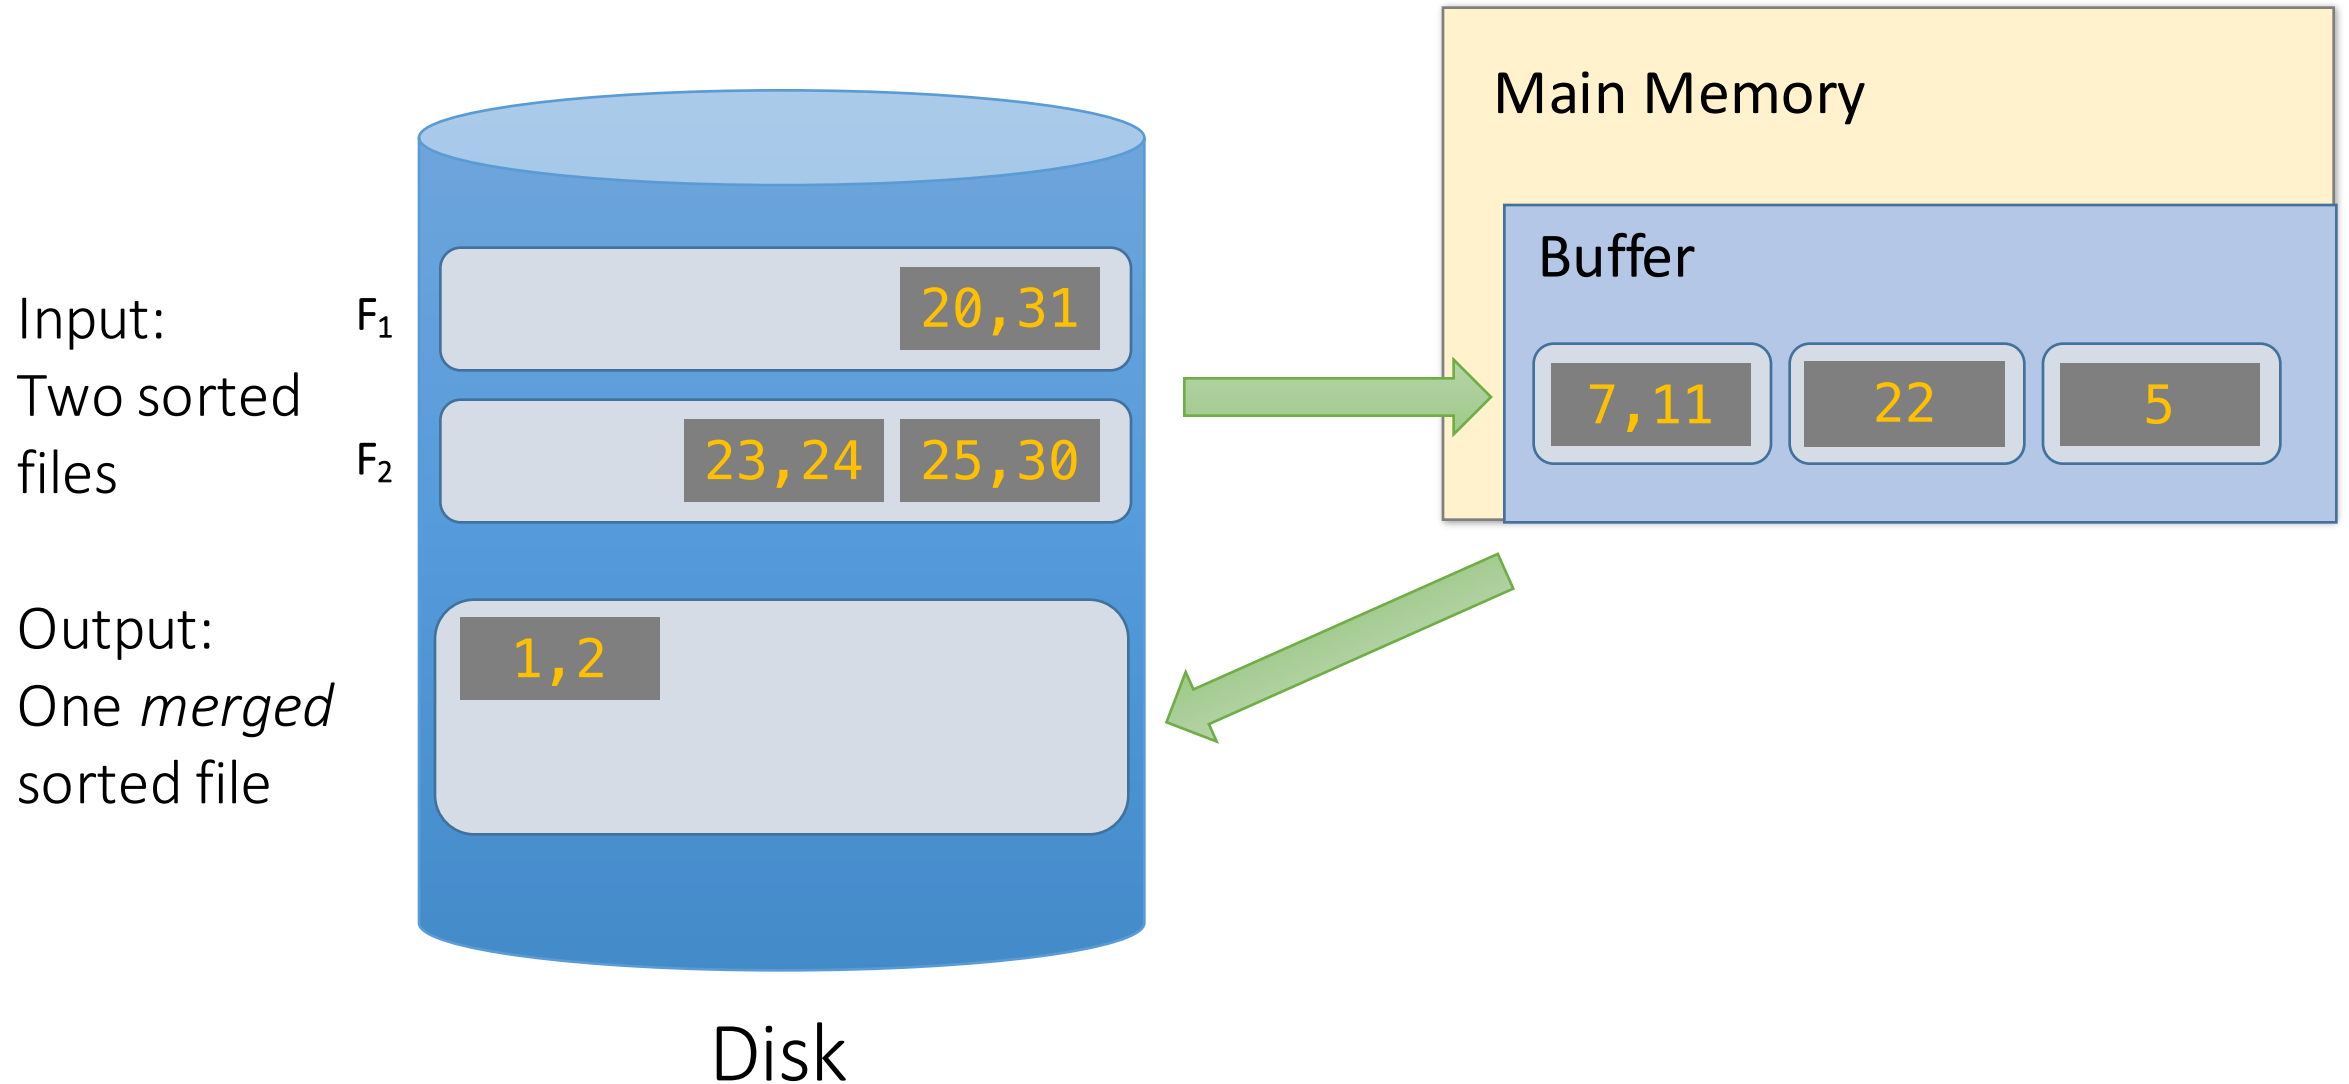
# External Merge Algorithm



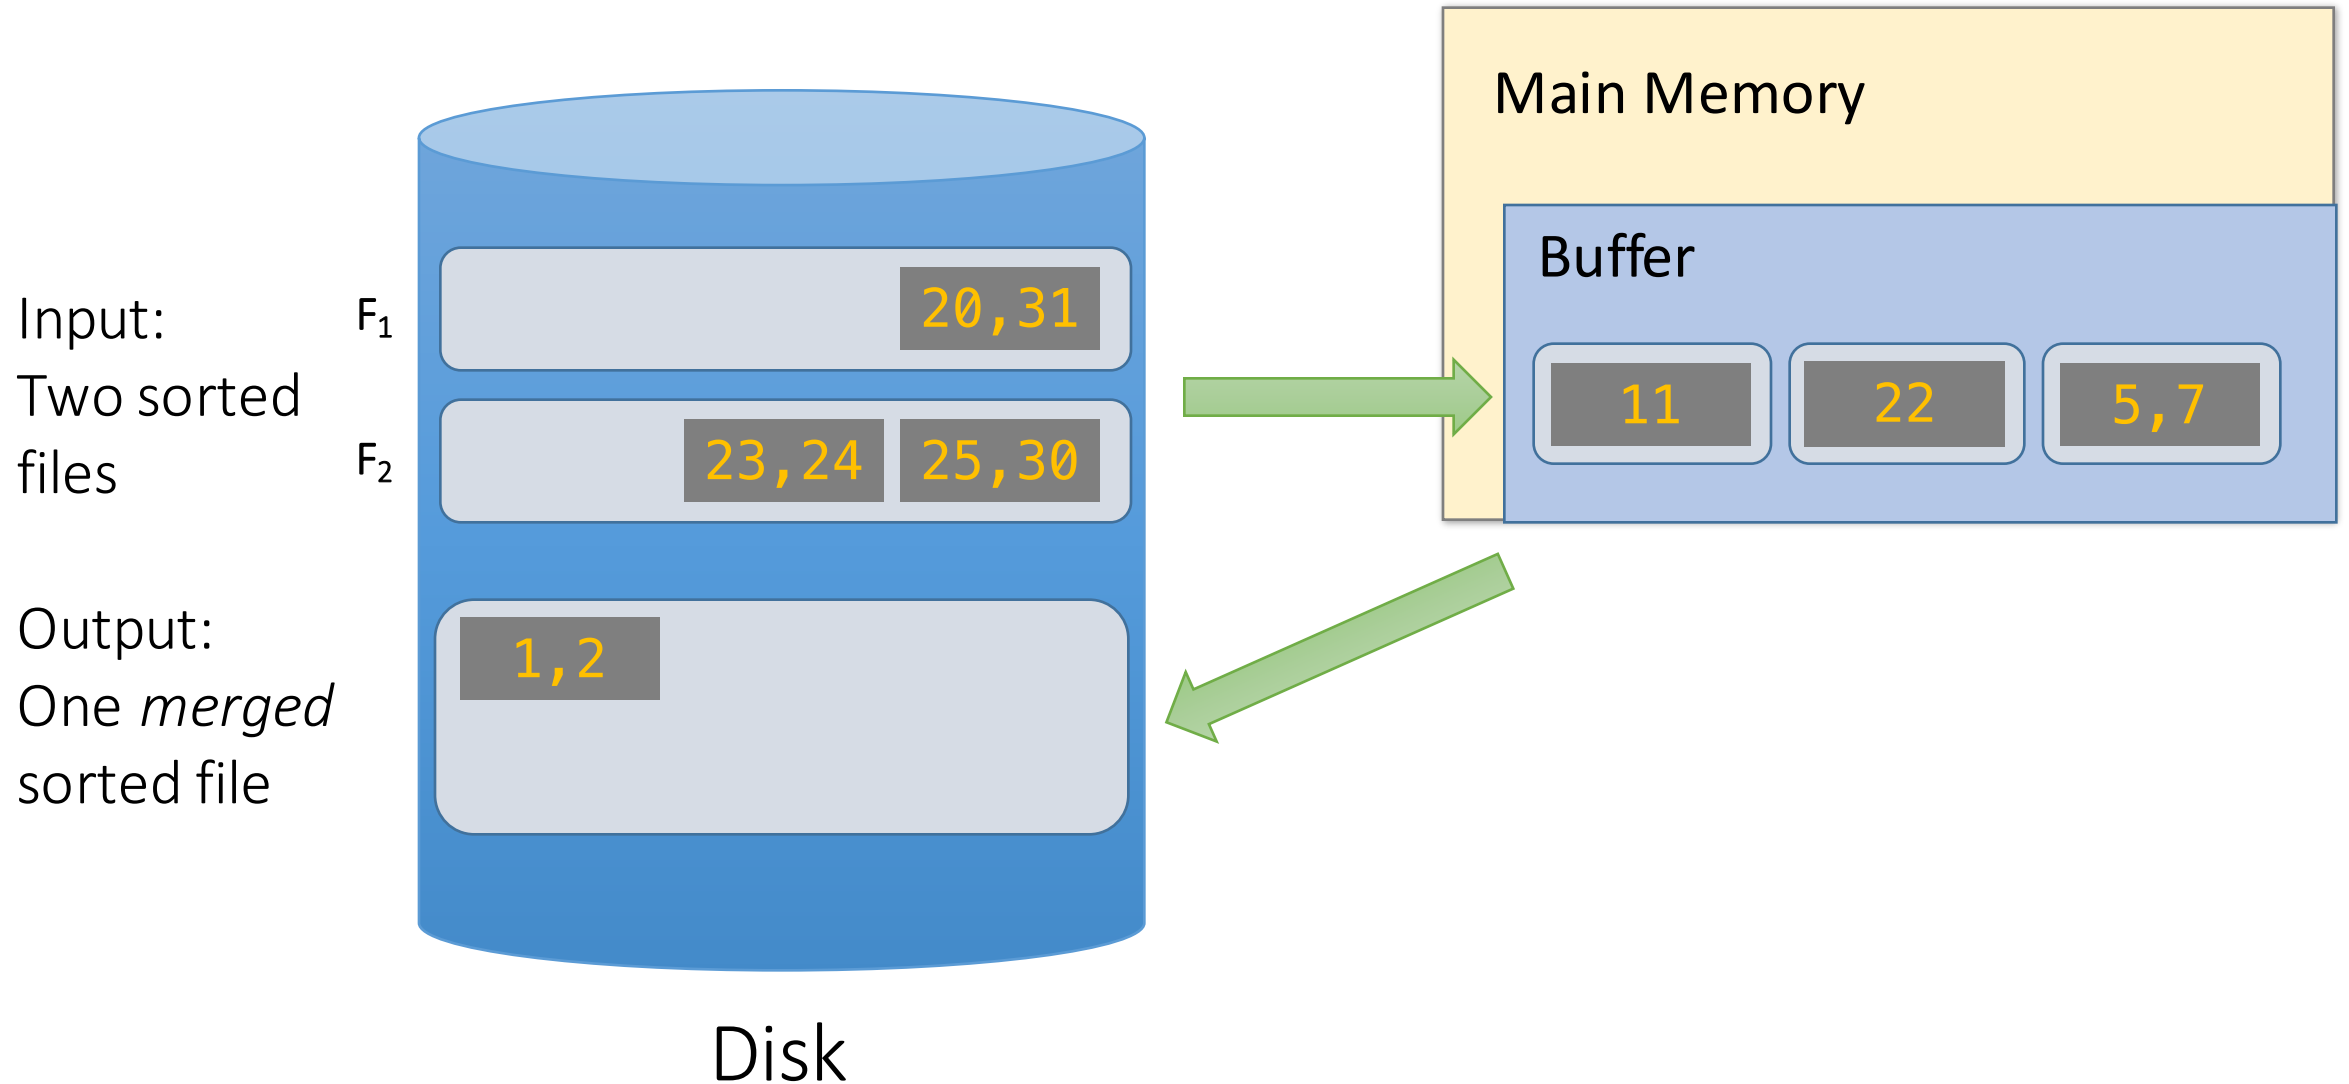
# External Merge Algorithm



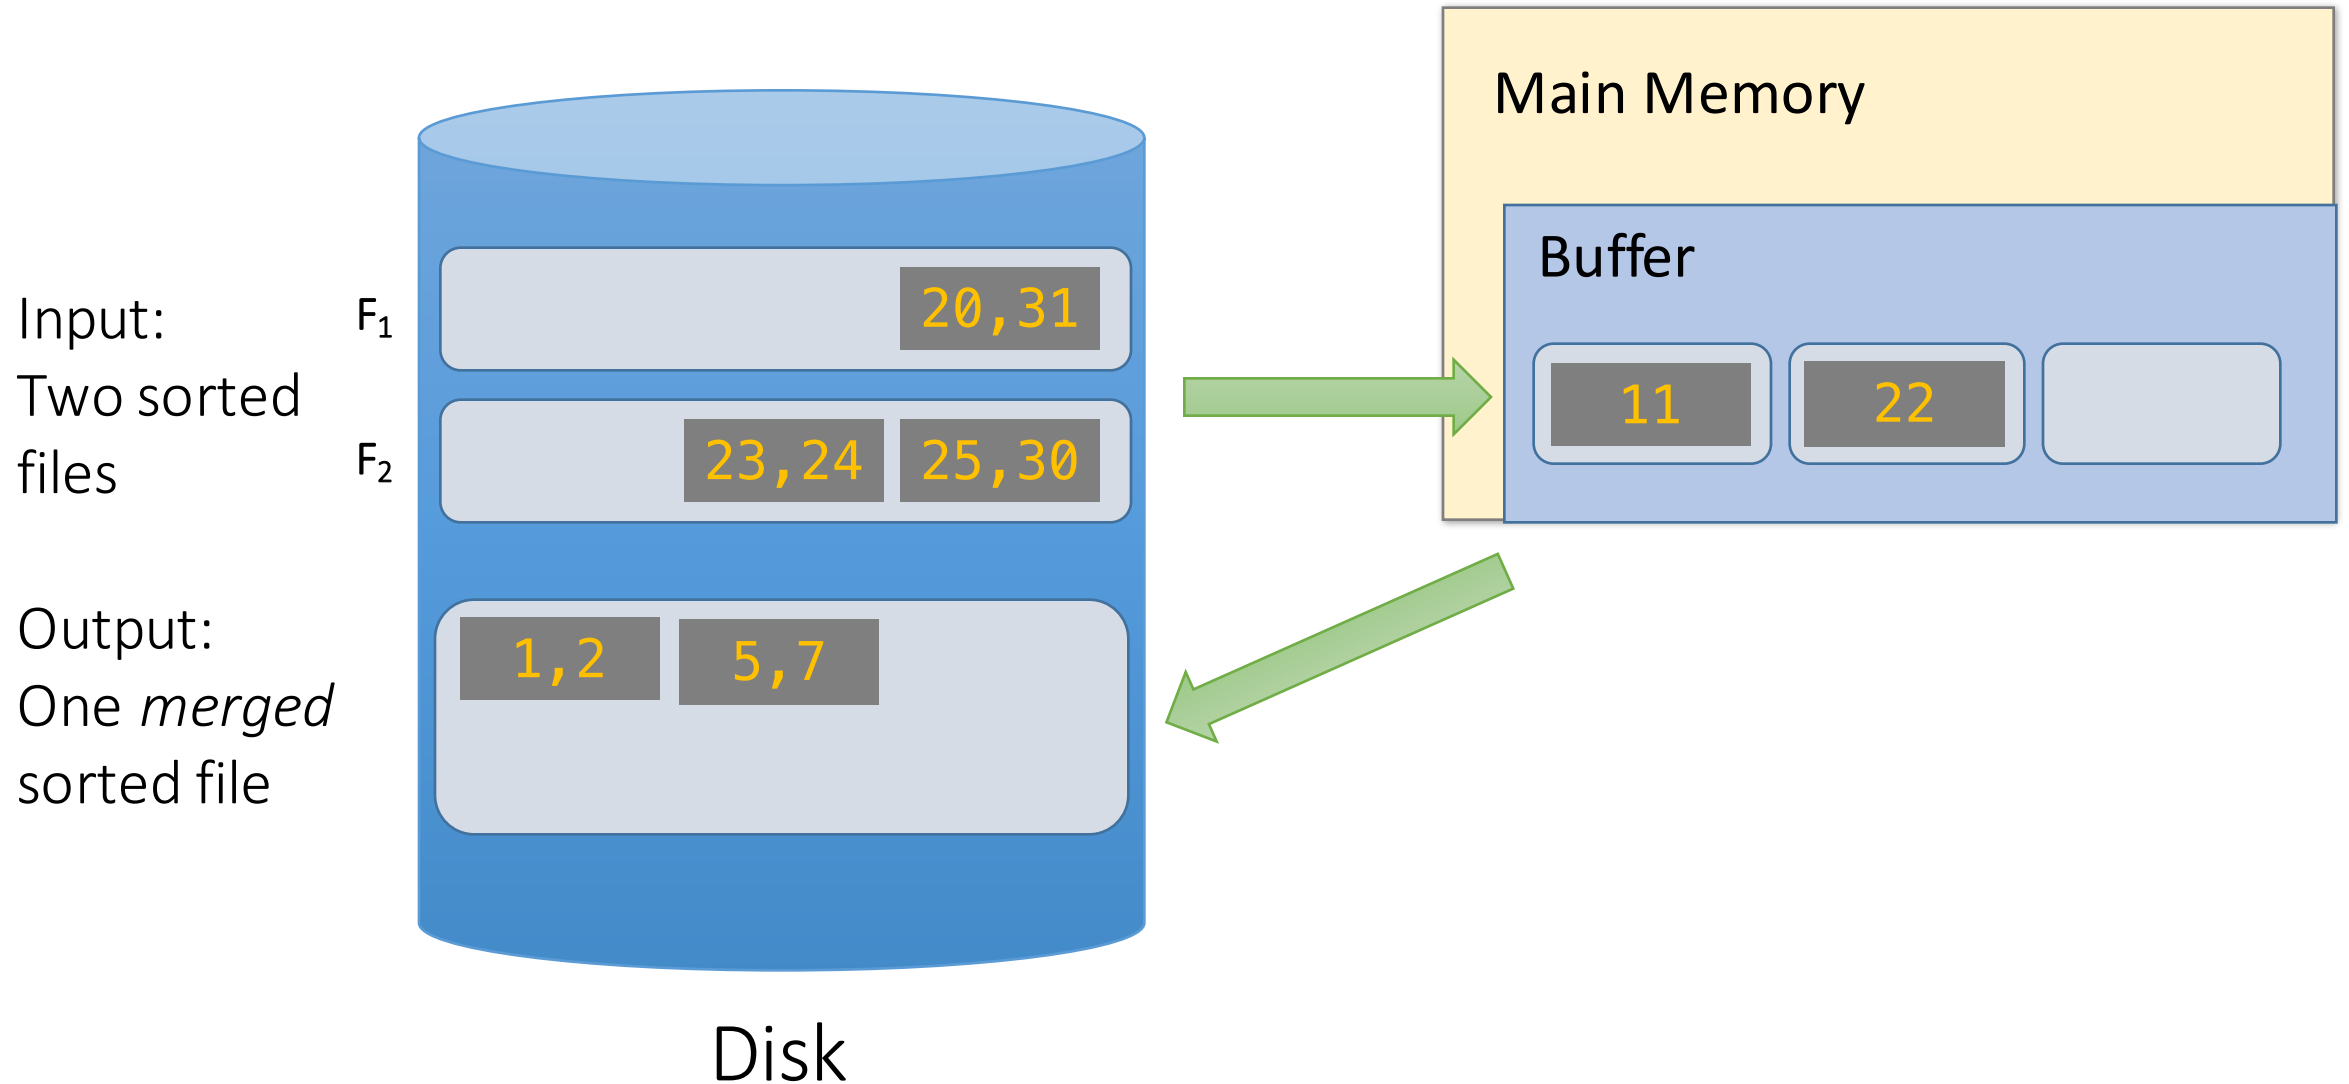
# External Merge Algorithm



# External Merge Algorithm

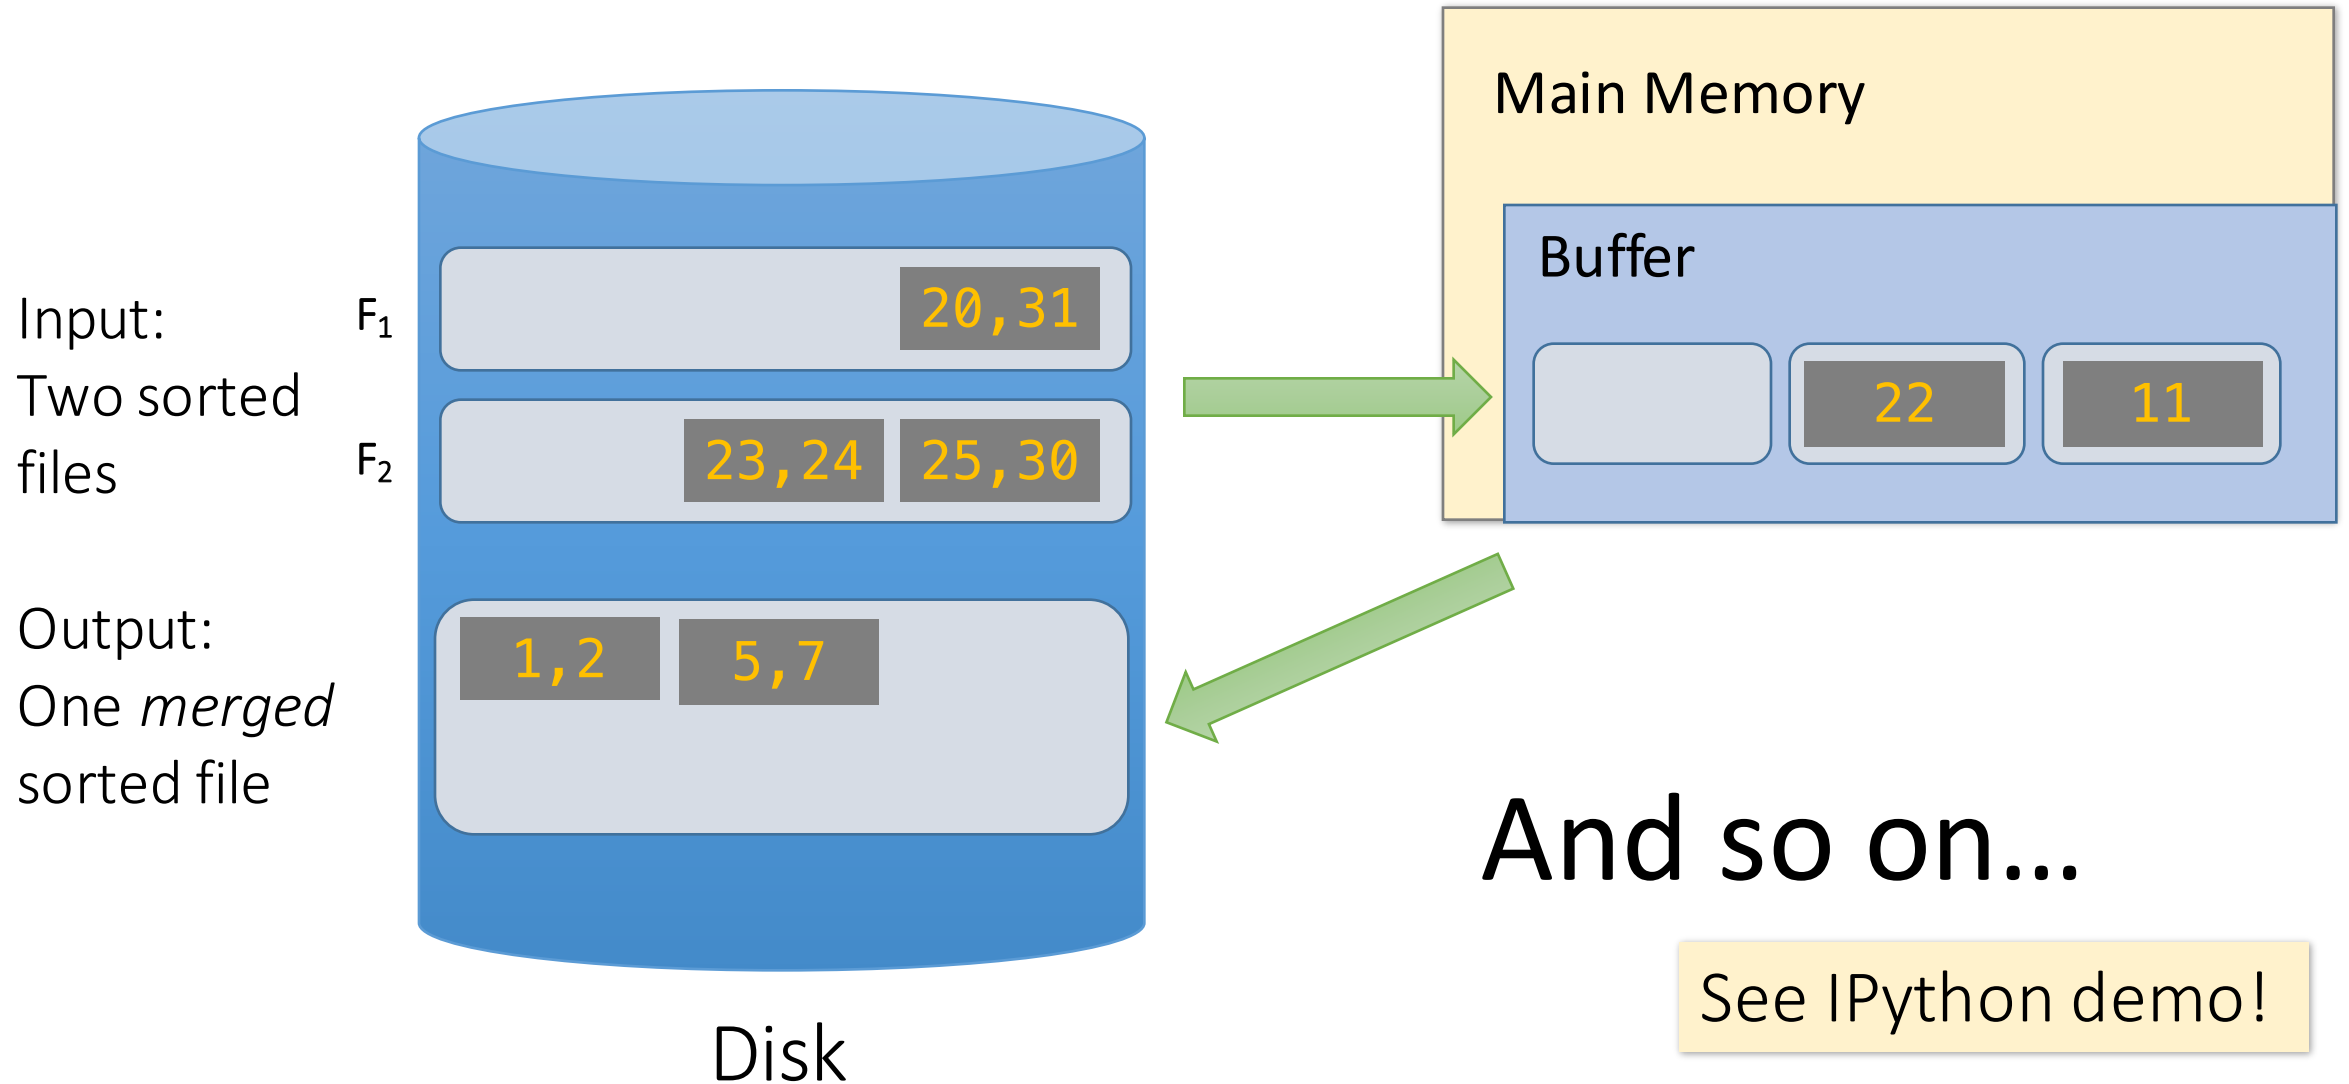


# External Merge Algorithm





# External Merge Algorithm



We can merge lists of **arbitrary length** with *only* 3 buffer pages.

If lists of size  $M$  and  $N$ , then

**Cost:**  $2(M+N)$  IOs

Each page is read once, written once

With  $B+1$  buffer pages, can merge  $B$  lists. How?